

Work Sample: Finding Where Agencies Post Their RFPs

Timebox: 2 hours. Partial completion is expected and fine.

This is the crux of the interview process. You do this, we talk once more to review, and I make a decision.

Context

ActiveParks identifies upcoming sales opportunities for vendors who sell big-ticket items to local government and school districts. Our specialty is catching projects at a very early stage, before anything is public.

We also track RFPs, which are the late-stage version of the same thing. When a city or a school district releases a request for proposal (RFP), our job is to spot it as soon as it goes up and get it to our customers fast. Being second is being late.

That only works if we know where to look. Every agency posts its solicitations somewhere, and there is no standard. Some publish on their own website. Many use a third-party procurement portal. Figuring out, per agency, which one it is and what the exact URL is, is the foundation. Get it wrong and we risk missing every RFP that agency ever posts.

The task

You'll get a CSV of 200 Indiana public agencies: 100 cities and 100 school districts. Four columns: agency name, state, agency type, and the agency's official website.

For each agency, find the page where they post RFPs, bids, and solicitations.

The point is not to find 200 URLs by hand. The point is to build an automated way to do it: an agent, a scraper, a script, a pipeline, whatever you think works. If you build something that works for 100 cities and 100 school districts, it should work for every city and every school district. How you approach that is most of what I'm evaluating.

The website column matters

You're given each agency's official website, and you should treat it as the identity of that agency, not just a convenience.

There is a Lawrence in Indiana, a Lawrence in Kansas, and a Lawrence in Massachusetts. A search for "Lawrence city bids" will happily return the wrong one, and that error is invisible

downstream. It looks like data. The website is how you prove the bid page you found belongs to the agency in front of you.

What these pages actually look like

Real examples from our own data.

Self-hosted, on the agency's own domain:

- <https://www.urbandale.org/bids.aspx?bidID=318>
- <https://thecityofniles.com/category/public-bids-2/>
- <https://www.oakdalemn.gov/1058/Park-Projects>

Official Third-party portals, where the agency has an account on a big aggregators's system:

- IonWave — <https://garlandtx.ionwave.net/PublicDetail.aspx?bidID=4047>
- PlanetBids — <https://vendors.planetbids.com/portal/15300/bo/bo-detail/144289>
- Bonfire — <https://dallascityhall.bonfirehub.com/opportunities/245284>
- CivicPlus — <https://sc-lancaster.civicplus.com/bids.aspx?bidID=41>
- Others you'll hit: DemandStar, Procureware, BidExpress, state-run portals.

Output

One CSV. These columns, in this order:

column	values
agency_name	copied verbatim from the input
state	copied verbatim
agency_type	copied verbatim
agency_website	copied verbatim
rfp_url	the page that actually lists current solicitations, not the homepage, not a generic "Departments" page
rfp_hosting	self-hosted / third-party / not-found
rfp_platform	vendor name if third-party, self-hosted, or not-found

Rules

- **Use a coding assistant.** Claude Code, Cursor, or Codex. I'd prefer Claude Code if you know it, but any of them is fine. This is not optional, and it's not a formality: a large part of what I'm hiring for is whether you can actually drive one of these tools on a real problem.
 - Beyond that, any language, any model, any approach. There is no intended solution.
 - **Not every agency has a clean answer.** Some post bids in a PDF, some only by email request, some sites are dead. A blank cell with a one-line reason in `notes` is worth more than a confident guess.
 - **Write down your assumptions.** If something was ambiguous and you made a call, say so. If you have questions, note them. I'd rather see a stated assumption than a silent one.
-

What to send back

1. **Your CSV**, whatever state it's in. If you didn't get to a CSV at all, that's a legitimate outcome. Send what you have.
2. **A short write-up.** Five sentences or five bullets is plenty:
 - What you got done in the two hours
 - Assumptions you made and questions you'd want answered
 - The failure mode that cost you the most time
 - What you'd do next with more time
3. **Your coding assistant session.** Instructions below. Messy is good. I want to see the actual working, including the parts that went sideways.

How to send your Claude Code session

Claude Code writes a full transcript of every session to a `.jsonl` file on your machine. To find the one you just created:

```
ls -t ~/.claude/projects/*/*.jsonl | head -5
```

The newest file is your session. Send it as an attachment:

```
cp "$(ls -t ~/.claude/projects/*/*.jsonl | head -1)"  
~/Desktop/my-session.jsonl
```

Two notes. These files can be large, so zip it if your mail client complains. And it's a raw log of everything in the session, so give it a look and strip anything private before you send it.

On Cursor: open the chat panel, use its export/share option, and send the resulting file or link.

If neither works, don't burn your two hours on it. Copy-paste the conversation into a text file, or send screenshots. I care about seeing your process, not the file format.

Timebox

Two hours from when you open the input file. Hard stop. Send whatever you have at that point.

Costs

This is not meant to cost you money. If you end up out of pocket, send me the receipt and I'll cover it. Realistically we're talking about a few dollars of tokens at most. If you find yourself heading toward real money, stop and tell me - that's a signal the approach has gone sideways, not a signal to keep paying.

If you want a third-party tool to do this properly, say so before you start the clock and I'll get you set up. We use a commercial scraping service (Scraping Bee) and I'm happy to give you access to it for this.